



Ha Noi University of Science and Technology

School of Information and Communications Technology

## **Parallel Programming Project**

**Topic:** Parallel Travelling Salesman Problem Solver  
Island-Model Genetic Algorithm with MPI

### **Group 19**

1. Dinh Quoc Bao - 20235478
  2. Chu Anh Duc - 20230081
  3. Vu Thuong Tin - 20230091
  4. Nguyen Xuan Khai - 20235508
- 

**Supervisor:** Dr. Nguyen Tuan Dung

June 24, 2026

## Abstract

This report describes a parallel solver for the Travelling Salesman Problem (TSP). The implementation uses an island-model Genetic Algorithm (GA) written in C++ and parallelized with MPI. Each MPI process is an independent island that evolves a population of tours. At fixed synchronization intervals, all islands cooperate by finding the global best tour with `MPI_Allreduce` and broadcasting the owner tour with `MPI_Bcast`. Other islands do not clone the whole tour; instead, they splice random contiguous segments of the global best into their own population through order crossover. This keeps the search diverse while still spreading useful sub-routes.

The repository also contains Python tools for visualization, validation, and plot generation. The Python code does not implement the TSP algorithm; it only drives the C++ executable or reads output files generated by the solver. The experimental results cover correctness validation, input-size selection, granularity/load balance, speedup with and without communication time, and a solution-quality comparison against a greedy baseline.

# Contents

|          |                                                        |           |
|----------|--------------------------------------------------------|-----------|
| <b>1</b> | <b>Project Overview</b>                                | <b>3</b>  |
| 1.1      | Problem Statement . . . . .                            | 3         |
| 1.2      | Notable Technical Points . . . . .                     | 3         |
| <b>2</b> | <b>Sequential Genetic Algorithm</b>                    | <b>4</b>  |
| 2.1      | Input, Output, and Representation . . . . .            | 4         |
| 2.2      | Fitness Evaluation . . . . .                           | 4         |
| 2.3      | Population Initialization . . . . .                    | 4         |
| 2.4      | One Generation of Evolution . . . . .                  | 4         |
| 2.5      | Selection, Crossover, and Mutation Operators . . . . . | 5         |
| 2.6      | Local Search . . . . .                                 | 5         |
| 2.7      | Greedy Baseline and Optional Seeding . . . . .         | 5         |
| <b>3</b> | <b>Mathematical Algorithm and EREW Analysis</b>        | <b>6</b>  |
| 3.1      | Mathematical Model . . . . .                           | 6         |
| 3.2      | $O(N)$ -Work Tour Evaluation on EREW . . . . .         | 6         |
| <b>4</b> | <b>Parallelization Design</b>                          | <b>7</b>  |
| 4.1      | Parallelism Level . . . . .                            | 7         |
| 4.2      | Decomposition Technique . . . . .                      | 7         |
| 4.3      | Mapping Technique . . . . .                            | 7         |
| 4.4      | Communication Strategy . . . . .                       | 7         |
| 4.5      | Migration Policy . . . . .                             | 8         |
| 4.6      | Load Balancing Considerations . . . . .                | 8         |
| 4.7      | Early Stopping . . . . .                               | 8         |
| <b>5</b> | <b>Parallel Algorithm Pseudocode</b>                   | <b>9</b>  |
| <b>6</b> | <b>Implementation and Instrumentation</b>              | <b>10</b> |
| 6.1      | Command-Line Interface and Metrics . . . . .           | 10        |
| 6.2      | Cluster Environment . . . . .                          | 10        |
| <b>7</b> | <b>Correctness Validation</b>                          | <b>11</b> |
| 7.1      | Permutation Validity . . . . .                         | 11        |
| 7.2      | Small-Instance Exact Check . . . . .                   | 11        |
| 7.3      | Parallel Communication Check . . . . .                 | 11        |
| 7.4      | Visualization Outputs . . . . .                        | 11        |
| <b>8</b> | <b>Experimental Results</b>                            | <b>13</b> |
| 8.1      | Experiment 1: Input Size Selection . . . . .           | 13        |
| 8.2      | Experiment 2: Granularity and Load Balance . . . . .   | 14        |
| 8.3      | Experiment 3: Speedup . . . . .                        | 14        |
| 8.4      | Experiment 4: Solution Quality . . . . .               | 15        |
| <b>9</b> | <b>Discussion</b>                                      | <b>17</b> |
| 9.1      | Why the Island Model Fits TSP . . . . .                | 17        |
| 9.2      | Communication Trade-Off . . . . .                      | 17        |
| 9.3      | Granularity . . . . .                                  | 17        |
| 9.4      | Network and Synchronization Effects . . . . .          | 17        |
| 9.5      | Quality Experiment Tooling . . . . .                   | 17        |
| 9.6      | Heterogeneous Node Effect . . . . .                    | 18        |

|                                            |           |
|--------------------------------------------|-----------|
| 9.7 Limitations . . . . .                  | 18        |
| <b>10 Line Count and Team Contribution</b> | <b>19</b> |
| 10.1 Line Count . . . . .                  | 19        |
| 10.2 Team Contribution . . . . .           | 19        |
| <b>11 Conclusion</b>                       | <b>19</b> |

# 1 Project Overview

## 1.1 Problem Statement

Given a set of  $N$  cities with two-dimensional coordinates, the Travelling Salesman Problem asks for a closed tour that visits every city exactly once and returns to the starting city while minimizing total Euclidean distance. For a tour  $T = (t_0, t_1, \dots, t_{N-1})$ , the objective value is:

$$L(T) = \sum_{i=0}^{N-1} d(t_i, t_{(i+1) \bmod N})$$

where  $d(a, b)$  is the Euclidean distance between city  $a$  and city  $b$ . TSP is NP-hard, so the project uses a heuristic search method instead of exact enumeration for large instances.

## 1.2 Notable Technical Points

Several design points are important for understanding why the implementation is structured this way:

- TSP is an NP-hard combinatorial optimization problem. Therefore, using a heuristic Genetic Algorithm is appropriate for large instances where exact enumeration is infeasible.
- A TSP tour is represented as a permutation of city indices. Standard one-point or two-point crossover can easily create duplicated or missing cities, so the solver uses permutation-preserving operators such as order crossover (OX).
- The parallelization is an island model, not a simple parallel `for`-loop. Each MPI rank runs a complete GA population and explores a different part of the search space.
- Communication uses MPI collective operations. The global best value and its owner are found with `MPI_Allreduce` using `MPI_MINLOC`, then the owner broadcasts the best tour with `MPI_Bcast`.
- The experiments analyze not only final running time but also speedup, load balance, communication overhead, and the effect of node heterogeneity on the cluster.

## 2 Sequential Genetic Algorithm

### 2.1 Input, Output, and Representation

The sequential GA is the core search procedure reused by each MPI island. The input is a city-coordinate file, where each non-comment line stores one city as an  $(x, y)$  point. After reading the file, cities are numbered  $0, 1, \dots, N - 1$  in file order.

An individual in the population is one candidate TSP tour. It is represented as a permutation of city indices:

$$[2, 0, 4, 1, 3]$$

The tour is closed implicitly, so the last city connects back to the first city. This representation is compact and makes validity easy to check: a legal tour must contain every index from  $0$  to  $N - 1$  exactly once. The output of the GA is the shortest tour found during the run and its total route length.

### 2.2 Fitness Evaluation

Before evolution starts, the solver precomputes a flat  $N \times N$  Euclidean distance matrix. Entry  $D[i \cdot n + j]$  stores the distance from city  $i$  to city  $j$ . This costs  $O(N^2)$  memory but makes later tour evaluation fast, because each edge lookup is  $O(1)$ .

For a tour  $T = (t_0, t_1, \dots, t_{N-1})$ , the fitness value is:

$$L(T) = \sum_{i=0}^{N-1} D[t_i, t_{(i+1) \bmod N}]$$

The solver minimizes  $L(T)$ , so a smaller value means a better individual. In code, this is implemented by `tour_length()` in `cpp/ga_core.hpp`.

### 2.3 Population Initialization

The default initialization creates `--pop` random permutations. Each permutation is generated by shuffling the city list, so the initial population is diverse and every individual is valid. If `--greedy-init` is enabled, only the first individual is replaced by the best parallel nearest-neighbor tour; the rest of the population remains random to avoid premature convergence.

### 2.4 One Generation of Evolution

Each generation creates a new population from the previous one. The implementation in `evolve_one_gen()` follows this sequence:

1. Sort individuals by tour length and copy the best individual unchanged into the next generation. This is elitism.
2. Repeatedly choose two parents by tournament selection.
3. Apply order crossover (OX) to produce one valid child permutation.
4. Mutate the child by random swap and segment reversal with probability `mut`.
5. Evaluate the child and append it to the new population.

After the new population reaches the original population size, it replaces the old population. This process is repeated until the generation limit is reached or, in the MPI solver, the global stagnation rule stops all islands.

## 2.5 Selection, Crossover, and Mutation Operators

The GA uses permutation-preserving operators:

- Tournament selection with tournament size  $k = 5$ . Five random individuals are sampled, and the shortest one is selected as a parent.
- Order crossover (OX), which copies a contiguous segment from one parent and fills the remaining positions using the order of the second parent. This preserves a useful sub-route while preventing duplicated or missing cities.
- Mutation first swaps two random cities and then reverses one random segment. These two small changes help the search escape local optima without destroying the entire tour.
- Elitism keeps one best individual per generation, guaranteeing that the best local solution on an island does not get lost by crossover or mutation.

## 2.6 Local Search

The optional memetic extension applies local search to the current best individual every `--twoopt` generations. This is not enabled in every experiment, but it is implemented for stronger local refinement. The local improvement consists of:

- 2-opt: remove two crossing edges and reconnect the tour by reversing a segment.
- Or-opt: move a short contiguous segment to a better insertion position.

## 2.7 Greedy Baseline and Optional Seeding

The current solver also contains a nearest-neighbor greedy construction in `cpp/tsp_island.cpp`. Each island can build a greedy tour from a different random start city. Then `MPI_Allreduce` with the `MINLOC` operator and `MPI_Bcast` select and share the best greedy tour across the cluster.

By default, the greedy tour is used only as a visualization/reference baseline when live output is enabled. If `--greedy-init` is passed, one individual in each island is seeded with this greedy tour while the rest of the population remains random. This gives the GA a good starting route without making all individuals identical.

### 3 Mathematical Algorithm and EREW Analysis

#### 3.1 Mathematical Model

Let the city set be  $V = \{0, 1, \dots, N-1\}$  and let a tour be a permutation  $T = (t_0, t_1, \dots, t_{N-1})$  of  $V$ . For Euclidean coordinates  $p_i = (x_i, y_i)$ , the edge weight is:

$$d(i, j) = \sqrt{(x_i - x_j)^2 + (y_i - y_j)^2}$$

The optimization problem solved by the GA is:

$$\min_{T \in S_N} L(T) = \min_{T \in S_N} \sum_{i=0}^{N-1} d(t_i, t_{(i+1) \bmod N})$$

where  $S_N$  is the set of all permutations of  $N$  cities. The exact search space has  $N!$  permutations, so the implementation uses randomized evolutionary search instead of enumeration.

#### 3.2 $O(N)$ -Work Tour Evaluation on EREW

The core mathematical kernel is evaluating one closed tour. The C++ function `tour_length()` performs this as a simple loop, so its sequential time is  $O(N)$ . The same computation is also valid under the EREW PRAM model (Exclusive Read, Exclusive Write) as an  $O(N)$ -work algorithm:

1. Processor  $P_i$  reads the two tour entries  $t_i$  and  $t_{(i+1) \bmod N}$  and writes one private value  $w_i = d(t_i, t_{(i+1) \bmod N})$ .
2. A tree reduction sums the private edge weights  $w_0, w_1, \dots, w_{N-1}$ .

Each edge-weight write targets a distinct memory cell, and the reduction can be scheduled so that every read and write in a step is exclusive. Therefore, the algorithm is EREW-safe. With  $N$  processors, the parallel time is  $O(\log N)$  for the reduction and the total work remains  $O(N)$ . In the implemented GA, each island evaluates many tours locally using the sequential  $O(N)$  loop; MPI parallelism is applied at the island/search level rather than inside a single tour evaluation.

## 4 Parallelization Design

### 4.1 Parallelism Level

The implementation uses coarse-grained task-level parallelism. Each MPI process owns a complete GA population and performs an independent stochastic search. The city data and distance matrix are replicated on every process.

The design also has an exploratory-search characteristic: different islands explore different regions of the search space due to different random seeds. Therefore, the decomposition is best described as a hybrid of task decomposition and exploratory decomposition.

### 4.2 Decomposition Technique

The selected decomposition is **exploratory decomposition with replicated input data**. Each island receives the same TSP instance but evolves a different population. This is suitable for a heuristic optimization problem because the main work is search, not deterministic matrix processing.

The alternatives were less suitable:

- Pure data decomposition is not natural because a single tour depends on the global order of all cities.
- Recursive decomposition is more appropriate for divide-and-conquer exact search, not for the selected GA.
- Speculative decomposition is possible for exact branch-and-bound but was not the focus of this project.

### 4.3 Mapping Technique

The mapping is a one-dimensional process-to-island assignment:

$$\text{rank } r \longrightarrow \text{island } r$$

Each rank stores:

- one local population of tours,
- one local vector of tour lengths,
- one local pseudo-random generator with seed `base + rank * 1000`,
- the replicated city coordinates and distance matrix.

For the cluster experiment, the hostfile uses a sequential round-robin list of hosts. This spreads ranks across four machines and supports up to 48 ranks, with 12 ranks per machine as the common usable rank budget. This cap follows the smallest machine in the cluster and avoids overloading stronger nodes with extra hyperthreads.

### 4.4 Communication Strategy

Communication is periodic and collective. Every `--sync` generations:

1. Each rank finds its local best individual.
2. `MPI_Allreduce` with `MPI_MINLOC` finds the globally best length and the owner rank.

3. The owner rank broadcasts the best tour using `MPI_Bcast`.
4. Non-owner ranks perform partial migration by combining a segment of the global best with local individuals.

This is a blocking collective communication strategy. The logical topology is global collective communication rather than a ring or master-slave topology. There is no permanent master rank during evolution; all ranks participate symmetrically. Rank 0 is used only for final reporting and file output.

## 4.5 Migration Policy

A naive island GA can broadcast the full best individual and clone it into every island. That is fast but dangerous: diversity collapses and all islands may converge to the same local optimum.

This implementation uses partial global-best migration:

- The global best tour is broadcast to all islands.
- Each non-owner island selects several worst individuals.
- For each replacement, it uses OX crossover between the global best and a random local mate.
- The child replaces one of the worst local individuals.

Only sub-routes are shared. Random cut points and local mates keep the resulting population different on each island.

## 4.6 Load Balancing Considerations

In the default mode, every rank uses the same population size, so compute work is approximately balanced. This is appropriate for machines with similar effective capacity or for a controlled experiment where ranks per machine are limited by the slowest machine.

The solver also contains an optional `--auto-balance` mode. It performs a small local benchmark on each process, gathers the measured times, and assigns a larger population to faster ranks. The allocation is clamped between 0.5x and 2x of the original population to avoid unstable decisions from noisy measurements.

## 4.7 Early Stopping

If `--patience` is enabled, the solver stops when the global best does not improve for a fixed number of generations. Because the same global best value is known to all ranks after each synchronization, every rank can make the same stop decision without an extra control message.

## 5 Parallel Algorithm Pseudocode

Listing 1: Island-model GA with partial global-best migration

```
Input: city file, generations G, population size P,  
      sync interval S, migrants M  
  
MPI_Init()  
rank <- MPI_Comm_rank()  
size <- MPI_Comm_size()  
  
coords <- read_cities(file)  
D <- distance_matrix(coords)  
rng <- seed + rank * 1000  
  
if greedy_init or live output enabled:  
  local_greedy <- nearest_neighbor_tour(D, random start)  
  (greedy_value, greedy_owner) <- MPI_Allreduce(MINLOC, local_greedy)  
  MPI_Bcast(best_greedy_tour, greedy_owner)  
  
pop <- P random tours  
len <- evaluate all tours  
if greedy_init:  
  pop[0] <- best_greedy_tour  
  len[0] <- greedy_value  
  
for g = 1 to G:  
  evolve_one_generation(pop, len, D)  
  
  if twoopt enabled and g % twoopt == 0:  
    polish local best with 2-opt and Or-opt  
  
  if S > 0 and size > 1 and g % S == 0:  
    local_best <- best tour on this rank  
    (global_value, owner) <- MPI_Allreduce(MINLOC, local_best)  
  
    if rank == owner:  
      buffer <- local_best_tour  
      MPI_Bcast(buffer, owner)  
  
    if rank != owner:  
      repeat M times:  
        worst <- worst local individual  
        mate <- random local individual  
        child <- OX(buffer, mate)  
        replace worst by child  
  
    update convergence-stop state  
    if stalled for patience generations:  
      break  
  
final_local_best <- best tour on this rank  
(global_best, owner) <- MPI_Allreduce(MINLOC, final_local_best)  
owner sends best tour to rank 0 if needed  
rank 0 writes route, history, stats  
MPI_Finalize()
```

## 6 Implementation and Instrumentation

### 6.1 Command-Line Interface and Metrics

The solver’s CLI configures GA parameters (`--gens`, `--pop`, `--patience`), MPI synchronization (`--sync`, `--migrants`), local search (`--twoopt`, `--greedy-init`), and data logging (`--out`, `--stats`, `--live`). Performance metrics recorded in the `--stats` CSV include communication time (`comm_s`), total elapsed time (`total_s`), estimated compute time (`compute_s`), and the overall execution makespan (`makespan_s`).

### 6.2 Cluster Environment

The cluster setup uses four machines and up to 48 MPI ranks. OpenMPI 5.0.9 is source-built and pinned at `/opt/openmpi-5.0.9` on all nodes to avoid runtime incompatibilities. The launch script uses:

Table 1: Cluster configuration used in the final experiments

| Node  | Hostname / user          | OS-visible CPU slots | Role              |
|-------|--------------------------|----------------------|-------------------|
| node1 | DESKTOP-9PH6P60 / bao    | 16                   | worker            |
| node2 | duck / duck              | 12                   | launcher + worker |
| node3 | LAPTOP-83S2JJ4P / acer   | 16                   | worker            |
| node4 | LAPTOP-B2AUBVCH / khainx | 16                   | worker            |

The final hostfile `cluster/hosts.seq48` contains 48 lines in round-robin order, so ranks are distributed as `node1,node2,node3,node4,node1,...`. Each machine receives 12 ranks. This mapping is one-dimensional: one MPI rank maps to one GA island.

- `--map-by seq`: sequential mapping from the hostfile.
- `--bind-to none`: avoids topology-dependent binding problems on heterogeneous machines.
- a pinned `prted` launch agent path for remote nodes.

## 7 Correctness Validation

### 7.1 Permutation Validity

The solver represents a route as a permutation of city IDs. The unit tests check that crossover and mutation preserve valid permutations. The validation script `python/validate_tour.py` independently verifies that a produced route:

1. has exactly  $N$  city indices,
2. contains every city from 0 to  $N - 1$  exactly once,
3. has a recomputed length matching the reported best length.

### 7.2 Small-Instance Exact Check

For  $N \leq 10$ , the validation script can brute-force all tours with city 0 fixed as the first city. This gives the true optimum and allows the GA result to be checked against the exact solution. The repository includes validation logs for small and larger instances in `results/`.

For large TSP instances, the GA is a heuristic method, so the validation does not prove global optimality. It confirms that the produced route is a valid TSP tour, that the reported length matches an independent recomputation, and that the solver can recover the exact optimum on the small brute-force test case.

### 7.3 Parallel Communication Check

The per-rank history files show that migration changes the behavior of non-owner islands after synchronization points. In one cluster run, rank 2 held the best tour before synchronization at generation 20. After the synchronization, weaker islands moved toward that better solution, but did not become identical. This matches the intended partial-migration behavior: useful route segments spread, while population diversity is preserved.

### 7.4 Visualization Outputs

The repository includes route and convergence visualizations. The current live viewer can read one JSONL stream per rank and show the greedy baseline before the GA evolution. These figures are useful for checking that the output route is a valid closed TSP tour and that the best length decreases over the generations.




Figure 1: Best TSP route visualization generated from the solver output.




Figure 2: Convergence history: best tour length over generations.

## 8 Experimental Results

### 8.1 Experiment 1: Input Size Selection

The requirement asks for an input-size experiment using the number of processes equal to the usable CPU slots of the cluster. The final cluster configuration uses 48 ranks, capped at 12 ranks per machine. The runtime is measured in two cases: total time including communication, and estimated compute time excluding communication.

The current `results/exp_size.csv` and `results/exp_size.png` come from a 48-rank size sweep at `gens=10500` and `sync=20`. The measured values are:

Table 2: Runtime versus input size with 48 MPI ranks, 10500 generations

| $N$  | Total time (s) | Communication (s) | Compute only (s) |
|------|----------------|-------------------|------------------|
| 1200 | 86.1346        | 51.1188           | 35.0158          |
| 1800 | 131.3597       | 65.6987           | 65.6610          |
| 2400 | 217.0933       | 100.3329          | 116.7604         |




Figure 3: Runtime sweep versus number of cities, with and without communication time.

This later sweep shows normal run-to-run variation on the heterogeneous cluster:  $N = 1800$  lies inside the 2–3 minute interval in this run, while  $N = 2400$  overshoots it. However, the dedicated N-star runs stored in the repository show that the selected final configuration also satisfies the target in repeated measurements. The final selection is:

Table 3: Dedicated  $N^*$  confirmation at 48 ranks, 10500 generations

| $N$  | Total time (s) | Evidence                                                                  |
|------|----------------|---------------------------------------------------------------------------|
| 2400 | 159.14         | repeated cluster run                                                      |
| 2400 | 154.27         | repeated cluster run                                                      |
| 2400 | 145.97         | <code>results/nstar_run.log</code> and <code>nstar_final_stats.csv</code> |

The selected input size is therefore:

$$N^* = 2400 \text{ cities, gens}^* = 10500, 48 \text{ ranks}$$

The official repeated run stored in `results/nstar_run.log` reports a makespan of **145.97 seconds**, which lies inside the required 120–180 second interval. Other repeated measurements at the same configuration were 159.14s and 154.27s, also inside the target range.

## 8.2 Experiment 2: Granularity and Load Balance

For the granularity experiment, the program is run with the selected input size  $N^* = 2400$ , `gens`=10500, and 48 ranks. Each rank is one island, so the granularity is one GA population per process. The required stacked bar chart shows compute time and communication time in the same column for each process.




Figure 4: Per-rank compute and communication time for the granularity experiment.

The latest `results/exp_gran.csv` gives a makespan of about 151.18s, with idle time ranging from 0.00s to 2.79s. The measured idle skew is therefore about 1.85%, which is below the 25% imbalance threshold. The system is considered load-balanced for this run. The main difference between ranks is not idle waiting but the split between compute time (65.34–129.84s) and communication time (21.33–83.05s), which is expected when ranks synchronize through blocking MPI collectives on heterogeneous machines.

## 8.3 Experiment 3: Speedup

The speedup experiment uses a fixed total workload and varies the number of processes. Following the requirement, this experiment uses  $2N^* = 4800$  cities with `gens`=400, `sync`=20, and a fixed total population of 480 individuals split over the active ranks. The final tested process count is 48, matching the usable rank budget of `hosts.seq48`; the next power-of-two value, 64, would exceed the prepared cluster allocation. The measured values are:

Table 4: Speedup experiment

| Procs | Total (s) | Comm (s) | Compute (s) | Speedup | No-comm speedup | Eff.   |
|-------|-----------|----------|-------------|---------|-----------------|--------|
| 1     | 16.3636   | 0.0000   | 16.3636     | 1.0000  | 1.0000          | 1.0000 |
| 2     | 9.6910    | 3.6071   | 6.0840      | 1.6885  | 2.6896          | 0.8443 |
| 4     | 4.7995    | 2.1130   | 2.6865      | 3.4094  | 6.0910          | 0.8524 |
| 8     | 3.3492    | 1.9690   | 1.3802      | 4.8858  | 11.8559         | 0.6107 |
| 16    | 2.8577    | 2.0346   | 0.8231      | 5.7261  | 19.8806         | 0.3579 |
| 32    | 3.6413    | 2.6965   | 0.9448      | 4.4939  | 17.3200         | 0.1404 |
| 48    | 4.0388    | 2.8897   | 1.1491      | 4.0516  | 14.2398         | 0.0844 |




Figure 5: Runtime and speedup as the number of processes changes.

The speedup improves up to 16 processes and then decreases. This means the parallel computation scales initially, but the communication and synchronization overhead becomes dominant at high process counts. The no-communication speedup is much higher, showing that the GA computation itself is parallelizable. The limiting factor is the cost of repeated collective communication on the real cluster network.

#### 8.4 Experiment 4: Solution Quality

The latest quality experiment compares the parallel greedy baseline against two GA modes at  $N = 1000$ , 48 ranks, `sync=20`, and normal `gens=1000`. The from-scratch GA is given 20 times more generations (`gens=20000`), while the greedy-seeded GA uses the normal generation budget and initializes one individual in each island from the best parallel nearest-neighbor tour.

Table 5: Solution-quality comparison from `results/exp_quality.csv`

| Mode            | Generations | Time (s) | Best length | Improvement vs. greedy |
|-----------------|-------------|----------|-------------|------------------------|
| Parallel greedy | 0           | 0.0264   | 2769.53     | 0.00%                  |
| GA from scratch | 20000       | 112.2263 | 3016.30     | -8.91%                 |
| GA greedy-seed  | 1000        | 4.5619   | 2594.70     | 6.31%                  |




Figure 6: Quality experiment: convergence curves and final best tour length.

This single latest run shows that greedy seeding gives the best final route under the normal generation budget, improving the greedy baseline by 6.31%. The from-scratch run is stochastic and, despite the larger generation budget, does not beat the greedy baseline in this measurement.

## 9 Discussion

### 9.1 Why the Island Model Fits TSP

The island model is a natural fit for TSP because the problem is a global combinatorial search. Multiple populations can explore different areas of the solution space without requiring communication at every generation. This reduces the cost compared with fine-grained parallel GA designs.

### 9.2 Communication Trade-Off

The `--sync` parameter controls the trade-off between cooperation and overhead. A small interval spreads good information quickly but increases communication cost and can reduce diversity. A large interval reduces overhead but makes the islands behave more like independent runs. The no-sharing mode `--sync 0` is useful as a baseline.

### 9.3 Granularity

The chosen granularity is coarse: one full GA population per process. This reduces communication frequency and makes the program easier to run on a heterogeneous cluster. The drawback is that very small TSP instances do not have enough computation per process, so communication dominates.

### 9.4 Network and Synchronization Effects

The project measured the real network to explain why communication time becomes important at high process counts. Launcher-to-worker latency was about 8–13 ms on average, while bandwidth through SSH transfer was about 18–21 MB/s. For small instances, the migration payload is tiny; for example, a 400-city tour contains only about 1.6 KB of integer data. Therefore, the main communication cost is latency and synchronization waiting, not raw bandwidth.

The final code supports `--sync 0`, which disables migration and acts as a no-sharing baseline. An additional measurement at  $N = 4800$ , 48 ranks, and 400 generations produced:

Table 6: Effect of synchronization interval at  $N = 4800$ , 48 ranks

| sync | Total (s) | Comm (s) | Compute (s) | Best length |
|------|-----------|----------|-------------|-------------|
| 0    | 38.29     | 0.55     | 37.75       | 174825.61   |
| 20   | 42.68     | 20.96    | 21.72       | 175298.77   |
| 100  | 44.57     | 21.23    | 23.34       | 175584.19   |
| 200  | 42.85     | 19.99    | 22.86       | 175368.18   |

Only `--sync 0` removes most communication time. Increasing the interval from 20 to 100 or 200 does not reduce measured communication time much because most of that time is synchronization waiting for slower ranks rather than pure data transfer. The best-length values in this table are single runs, so they are not used to claim that one synchronization interval gives better solution quality than another.

### 9.5 Quality Experiment Tooling

The current `python/experiments.py` also provides a `quality` experiment. It compares three arms: parallel nearest-neighbor greedy, GA-from-scratch, and GA with `--greedy-init`. This experiment is useful for demonstrating solution quality, but the required grading experiments remain input-size selection, granularity, and speedup.

## 9.6 Heterogeneous Node Effect

The cluster is not homogeneous. A follow-up check removed node3 and compared the original 48-rank run against two three-node configurations. Using too many hyperthreads on the remaining nodes was slower, while a capped 36-rank run without node3 was slightly faster than the 48-rank baseline. This indicates that node3 behaves as a mild straggler, most likely due to memory pressure and general system heterogeneity. The report therefore avoids claiming that RAM is the only cause.

## 9.7 Limitations

The current implementation is a heuristic solver. It does not guarantee the global optimum for large TSP instances. The exact check is only feasible for small  $N$ . The selected  $N^* = 2400$  configuration satisfies the 2–3 minute runtime target, but the result is still affected by normal run-to-run variation on a heterogeneous laptop cluster.

## 10 Line Count and Team Contribution

### 10.1 Line Count

The professor requires at least 250 lines of code per member. The clean source-code count, excluding large generated data and binary report files, is summarized below:

Table 7: Lines of code by module

| Module          | Main files                                                                                                      | LOC  |
|-----------------|-----------------------------------------------------------------------------------------------------------------|------|
| C++ solver      | <code>tsp_island.cpp</code> , <code>ga_core.hpp</code> , <code>local_search.hpp</code> , <code>tests</code>     | 650  |
| Python tooling  | <code>experiments.py</code> , <code>benchmark.py</code> , <code>visualize.py</code> , <code>live_view.py</code> | 785  |
| Cluster scripts | installation, SSH setup, synchronization, launcher scripts                                                      | 265  |
| Total           | source code and scripts                                                                                         | 1700 |

For a four-member group, this is approximately 425 lines per member, above the 250-line threshold.

### 10.2 Team Contribution

The project combines algorithm implementation, MPI parallelization, cluster operation, experiments, and report writing. The following table records the report-level responsibility split used for submission. If the team’s internal work log differs, these entries should be adjusted before final submission.

Table 8: Member responsibility summary

| Member           | Main contribution                                                                                                                           |
|------------------|---------------------------------------------------------------------------------------------------------------------------------------------|
| Dinh Quoc Bao    | C++ GA core, tour representation, fitness evaluation, OX crossover, mutation, local-search support, and correctness checks.                 |
| Chu Anh Duc      | MPI island solver, <code>Allreduce</code> / <code>Bcast</code> migration logic, timing/statistics output, benchmark and experiment scripts. |
| Vu Thuong Tin    | Cluster preparation, OpenMPI/SSH setup, hostfile design, code synchronization, and multi-node execution support.                            |
| Nguyen Xuan Khai | Experiment runs, result figures, validation logs, report organization, visualization, and final analysis.                                   |

## 11 Conclusion

This project implements a complete MPI-based parallel TSP solver using an island-model Genetic Algorithm. The parallelism is coarse-grained and task/exploratory in nature. Each MPI rank owns an island and communicates only periodically through blocking MPI collectives. The mapping is one-dimensional: rank to island. The communication topology is global collective communication using `MPI Allreduce` and `MPI Bcast`.

Correctness is checked through permutation tests, independent route-length validation, and brute-force validation for small instances. The experimental results show balanced work distribution and meaningful speedup up to a moderate number of processes. At high process counts, communication overhead dominates, which is visible in the difference between total speedup and no-communication speedup.

## Appendix: Reproducible Commands

```
# Build
cd cpp
make
make test
cd ..

# Local MPI run
mpirun --oversubscribe -np 4 ./cpp/tsp_island \
  data/cities_50.txt --gens 500 --sync 20

# Cluster run
bash cluster/run_cluster.sh cluster/hosts.seq48 48 \
  ./cpp/tsp_island data/cities_2400.txt \
  --gens 10500 --sync 20 --stats results/nstar_final_stats.csv

# Required experiment figures
python3 python/experiments.py size \
  --procs 48 --sizes 1200 1800 2400 \
  --gens 10500 --sync 20 --hostfile cluster/hosts.seq48

python3 python/experiments.py gran \
  --procs 48 --size 2400 --gens 10500 --sync 20 \
  --hostfile cluster/hosts.seq48

python3 python/experiments.py speedup \
  --procs 1 2 4 8 16 32 48 --size 4800 \
  --gens 400 --sync 20 --hostfile cluster/hosts.seq48

# Optional quality comparison
python3 python/experiments.py quality \
  --procs 48 --size 1000 --gens 1000 --scratch-mult 20 \
  --sync 20 --hostfile cluster/hosts.seq48

# Live demo GIFs
python3 python/live_view.py run data/cities_100.txt \
  --islands 4 --gens 800 --sync 30 --save results/live100.gif

python3 python/live_view.py run data/cities_100.txt \
  --islands 4 --gens 800 --sync 30 --greedy-init \
  --save results/live100_greedy.gif
```